

Exploring word2vec vector space

JULIA BAZIŃSKA



„a word is characterized by the company it keeps”

John Rupert Firth

Word similarity

$P(a|c)$ – probability that word a appears in the neighbourhood of word c .

Words a and b have
similar meanings



$P(c|a) \approx P(c|b)$
for every word c

But then we would need to know $P(x/y)$ for every pair x ,
 $y \dots$

Pointwise mutual information

$$PMI(x, y) = \log \left(\frac{P(x \wedge y)}{P(x)P(y)} \right) = \log \left(\frac{P(x | y)}{P(x)} \right)$$

How much more probable are words x , y to occur together than at random?

Vector approximation

For words x and y let's find vectors v_x and v_y satisfying:

$$PMI(x, y) = \overrightarrow{v_x} \cdot \overrightarrow{v_y}$$

Back to word similarity

$$P(c|a) \approx P(c|b) \quad \text{for every word } c.$$

$$PMI(a, c) \approx PMI(b, c)$$

$$\vec{v}_a \cdot \vec{v}_c \approx \vec{v}_b \cdot \vec{v}_c$$

$$\vec{v}_c \cdot (\vec{v}_a - \vec{v}_b) \approx 0$$

$$\vec{v}_a \approx \vec{v}_b$$

Cosine distance

- Similar words have similar vector values
- We can use cosine distance to measure similarity:

$$\text{dist}(a, b) = \frac{\vec{v}_a \cdot \vec{v}_b}{|\vec{v}_a| |\vec{v}_b|}$$

```
find_most_similar("blue", 10)
```

```
0  
blue      1.000000  
red       0.890182  
black     0.864808  
pink      0.845264  
green     0.834646  
yellow    0.832033  
purple    0.829353  
white     0.822612  
orange    0.811403  
bright    0.799914  
dtype: float64
```

```
find_most_similar("dance", 10)
```

```
0  
dance     1.000000  
dancing   0.906835  
singing   0.871643  
dances    0.853638  
music     0.853078  
musical   0.839160  
dancers   0.813865  
hop       0.797830  
singers   0.788027  
pop       0.787464  
dtype: float64
```



```
find_most_similar("europe", 10)
```

```
0
europe          1.000000
european        0.881251
asia            0.834660
world           0.826442
countries       0.819641
britain         0.816004
continent       0.798280
america         0.794502
germany         0.791652
country         0.790833
```

```
find_most_similar("vinci", 10)
```

```
0
vinci           1.000000
leonardo        0.739511
botticelli      0.672160
michelangelo    0.661199
caravaggio      0.658187
vaio            0.641079
andrea          0.631471
giovanni        0.628268
vita            0.627567
francesca       0.626247
```

X is to Y as A is to...?

The nearest vector to $v_y - v_x + v_a$ is v_b with the highest value of $v_b \cdot (v_y - v_x + v_a)$

$$v_b \cdot (v_y - v_x + v_a) = v_b \cdot v_y + v_b \cdot v_x - v_b \cdot v_a$$

```
riddle("warsaw", "poland", "moscow")
```

```
0
```

russia	0.955646
ukraine	0.871337
russian	0.849050
poland	0.846253
republic	0.842461

```
riddle("good", "bad", "up")
```

```
0
```

down	0.953188
up	0.920470
falling	0.912033
out	0.893837
dropping	0.876727

```
riddle("hope", "disappointment", "peace")
```

0

disagreement	0.764241
disappointment	0.751991
underlined	0.742288
renewed	0.714227
underscored	0.712070

```
riddle("country", "language", "britain")
```

0

english	0.780426
language	0.777146
translation	0.762685
text	0.745039
pronunciation	0.729750


```
riddle("science", "einstein", "painting")
```

```
0
```

```
matisse      1.117469
```

```
picasso      1.083420
```

```
duchamp      1.055209
```

```
rembrandt    1.027414
```

```
titian        1.001785
```

```
dtype: float64
```

```
riddle("astronomy", "copernicus", "philosophy")
```

```
0
```

```
nietzsche    0.897995
```

```
copernicus    0.886114
```

```
freud         0.872768
```

```
hegel         0.865735
```

```
kierkegaard   0.843774
```

Projection on a difference axis

To measure whether v_x is associated more with v_a or v_b we can calculate

$$v_x \cdot (v_a - v_b)$$

- Name gender
- Good – bad vs up – down
- Interactive viz: lamyowce.github.io/word2viz

Vector rejection

$$v'_a = v_a - \frac{v_a \cdot v_b}{\underbrace{v_b \cdot v_b}} \cdot v_b \Rightarrow v'_a \perp v_b$$

projection of a on b

Vector rejection

- Can help with polysemy problems!

$$v'_{rock} = v_{rock} - \frac{v_{rock} \cdot v_{music}}{v_{music} \cdot v_{music}} \cdot v_{music}$$


```
find_most_similar('rock', dfn, 10)
```

```
0
```

rock	1.000000
band	0.694770
pop	0.674067
punk	0.661065
bands	0.645949
'n'	0.624487
rocks	0.616312
album	0.613320
albums	0.600090
music	0.599088

```
dfn.dot(reject('rock', ['music'], dfn)).sort_values(ascending = False).head(10)
```

```
0
rock          0.641093
rocks          0.498883
outcropping    0.444147
limestone      0.431764
outcrops       0.419863
outcrop        0.409561
cliffs         0.404085
fraggle        0.400296
rockers        0.399462
granite        0.399351
```

```
find_most_similar('python', dfn, 10)
```

```
0
```

python	1.000000
monty	0.689272
spamalot	0.561178
cleese	0.545438
php	0.525527
pythons	0.507684
perl	0.499981
scripting	0.485102
skit	0.475383
reticulatus	0.470973

```
dtype: float64
```

```
dfn.dot(reject('python', ['monty'], dfn)).sort_values(ascending = False).head(10)
```

```
0  
python          0.524904  
php             0.414212  
scripting       0.398441  
java            0.336685  
server-side     0.334762  
pythons         0.323510  
perl            0.322435  
javascript      0.317562  
c++             0.316876  
bindings        0.316854
```

```
dfn.dot(reject('python', ['monty', 'php'], dfn)).sort_values(ascending = False).head(10)
```

```
0  
python          0.307225  
pythons         0.263030  
reticulated     0.236272  
crocodile       0.232522  
snake           0.230709  
monkey          0.227266  
burmese         0.225066  
lizard          0.219475  
turtles         0.216014  
tortoise        0.214143
```



```
find_most_similar("polish", dfn)
```

```
0  
polish      1.000000  
lithuanian  0.725827  
hungarian   0.701218  
poland      0.694825  
slovak      0.667826  
dtype: float64
```

```
dfn.dot(reject('polish', ['lithuanian'], dfn)).sort_values(ascending = False).head(10)
```

```
0  
polish      0.473175  
poland      0.383410  
warsaw      0.344952  
german      0.286086  
jerzy       0.267660  
walesa      0.263369  
pope        0.262397  
krakow      0.260597  
lech        0.253146  
iraqi       0.251096
```

Removing gender from vectors

- We can reject $v_{she} - v_{he}$ from all vectors in data

```
find_most_similar('jane', dfn, 10)
```

0	
jane	1.000000
elizabeth	0.608416
anne	0.601765
mary	0.572299
sarah	0.570372
eliza	0.558833
alice	0.555983
helen	0.553562
ellen	0.553515
fonda	0.537337

```
find_most_similar('mother', dfn, 10)
```

0	
mother	1.000000
daughter	0.864802
wife	0.856802
grandmother	0.837379
husband	0.805565
sister	0.802924
father	0.793677
her	0.783749
daughters	0.758976
woman	0.757987



```
find_most_similar('jane', ungended, 10)
```

0	
jane	1.000000
elizabeth	0.556611
anne	0.545300
john	0.540406
bruce	0.537315
william	0.527928
henry	0.525371
eyre	0.515861
mary	0.515024
sarah	0.506965



```
find_most_similar('mother', ungended, 10)
```

0	
mother	1.000000
father	0.888926
daughter	0.842649
wife	0.837058
grandmother	0.811116
husband	0.804543
son	0.787408
sister	0.759808
brother	0.757006
her	0.738656

Vector quality

- Corpus size
- Vector dimensions
- Vector deriving algorithm

The background is a dark navy blue. In the top-left corner, there are several parallel teal lines forming a right-angled corner. In the bottom-left corner, there are several parallel teal lines forming a right-angled corner. In the bottom-right corner, there are several parallel teal lines forming a right-angled corner.

Thank you

Julia Bazińska
julia@bazinski.com
GH: lamyiowce